AWS Deployment

Insight Data Science

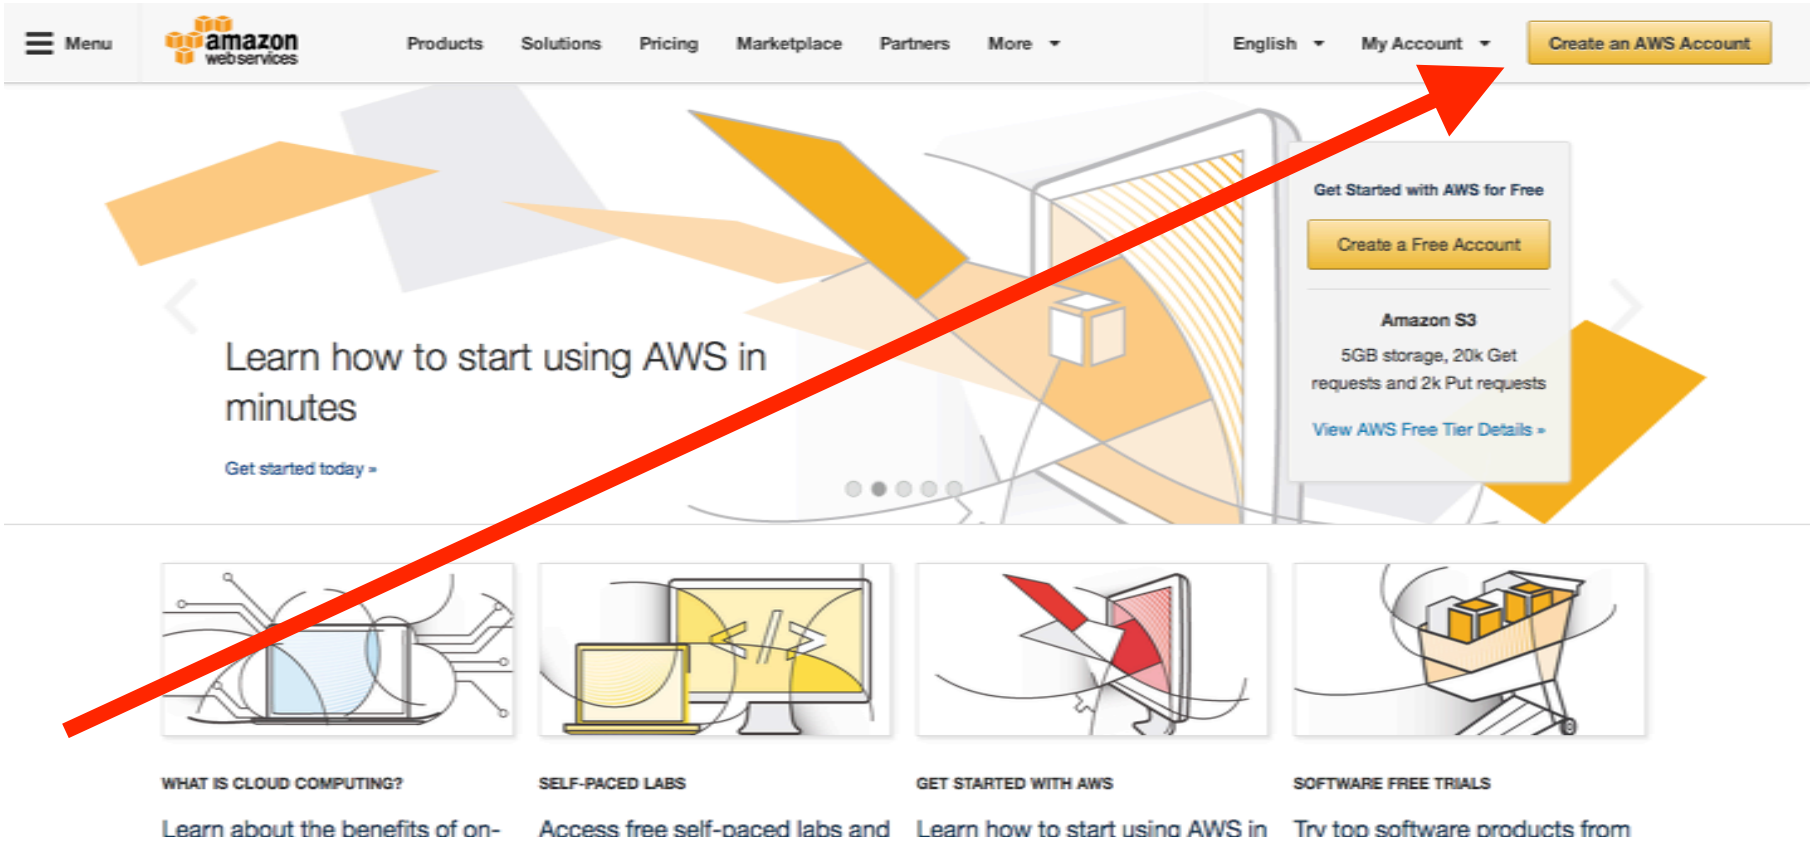

Sign up at aws.amazon.com

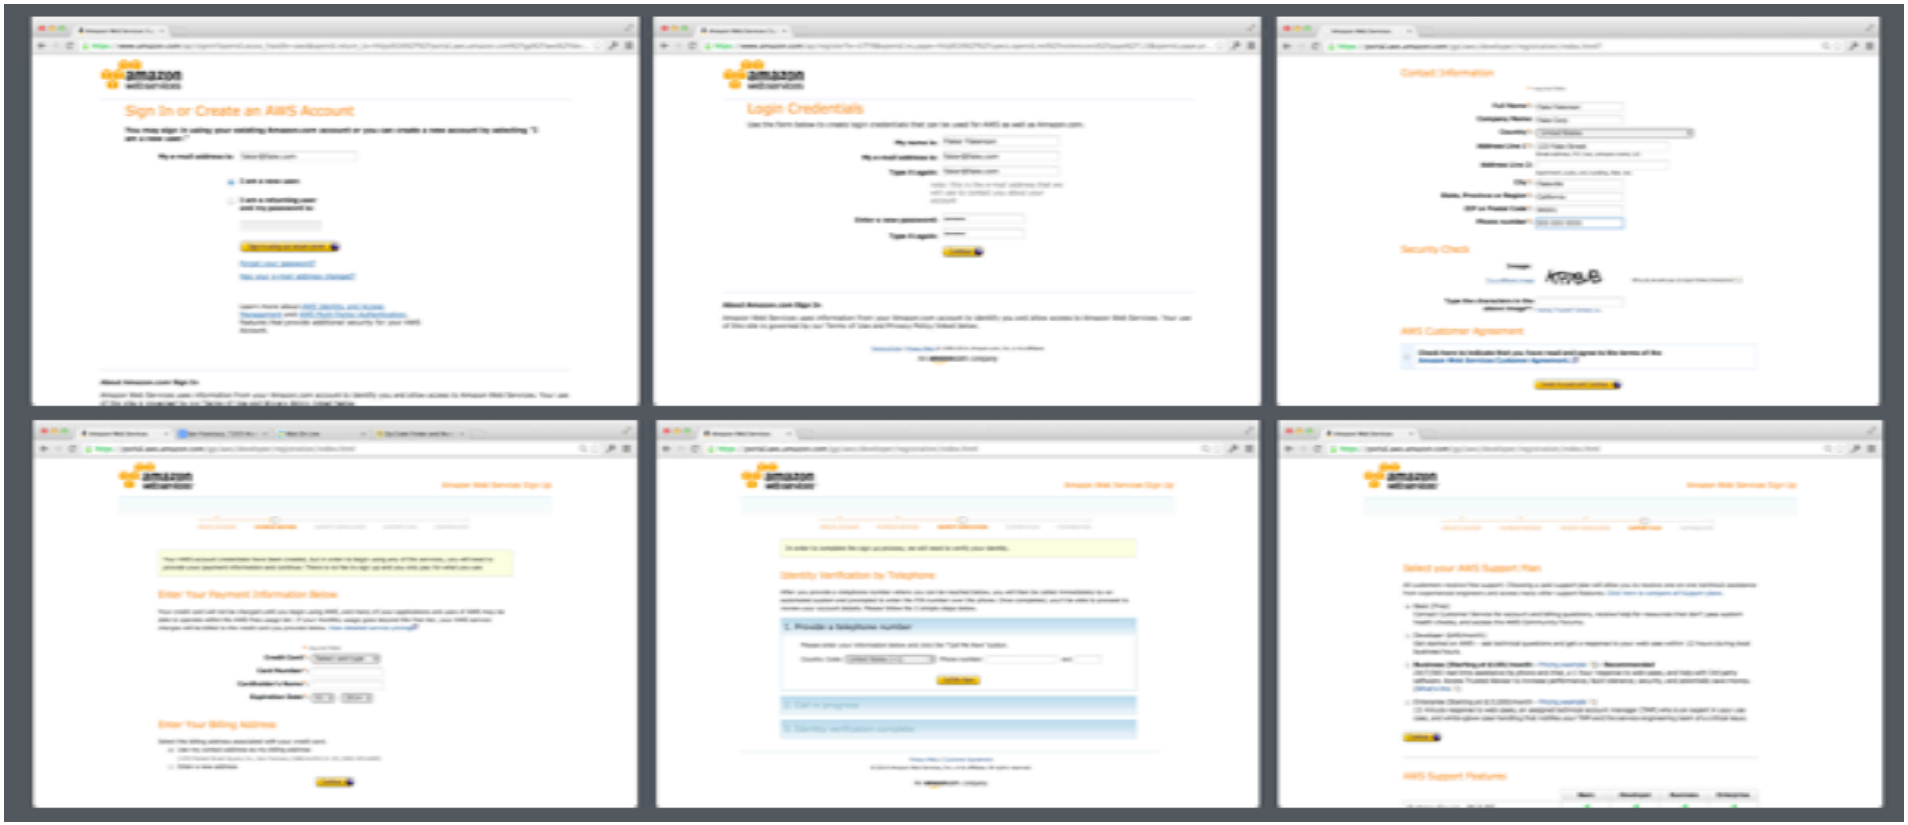

You'll need to complete a few forms...

Two things of note

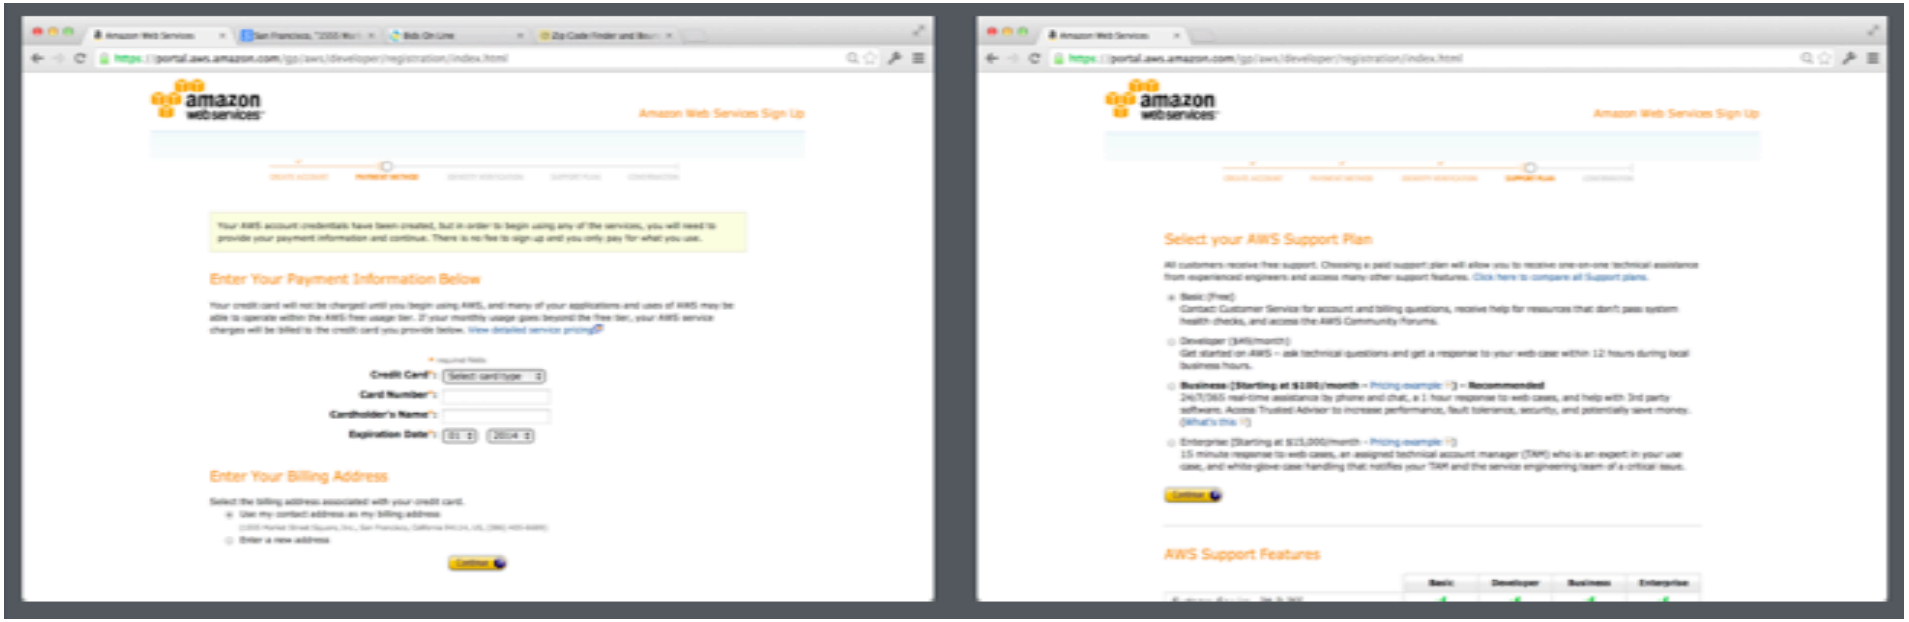

Credit card needed

Select Basic Support Plan - it's free!

You now have an AWS account!!!

Let’s log in


Amazon Web Services, Cl
aws.amazon.com
amazon web services
Sign Up
My Account / Console
English
AWS Products & Solutions
AWS Management Console
My Account
Billing & Cost Management
Security Credentials
Launch Software into VPCs with 1-Click
Now it's even easier to launch AWS Marketplace software into VPCs.
Get started »
VPC
Featuring
riverbed
CITRIX
SOPHOS
Silver Peak
Check Point
Get Started for Free »
Launch virtual machines and apps in minutes.
AWS
AWS
https://console.aws.amazon.com

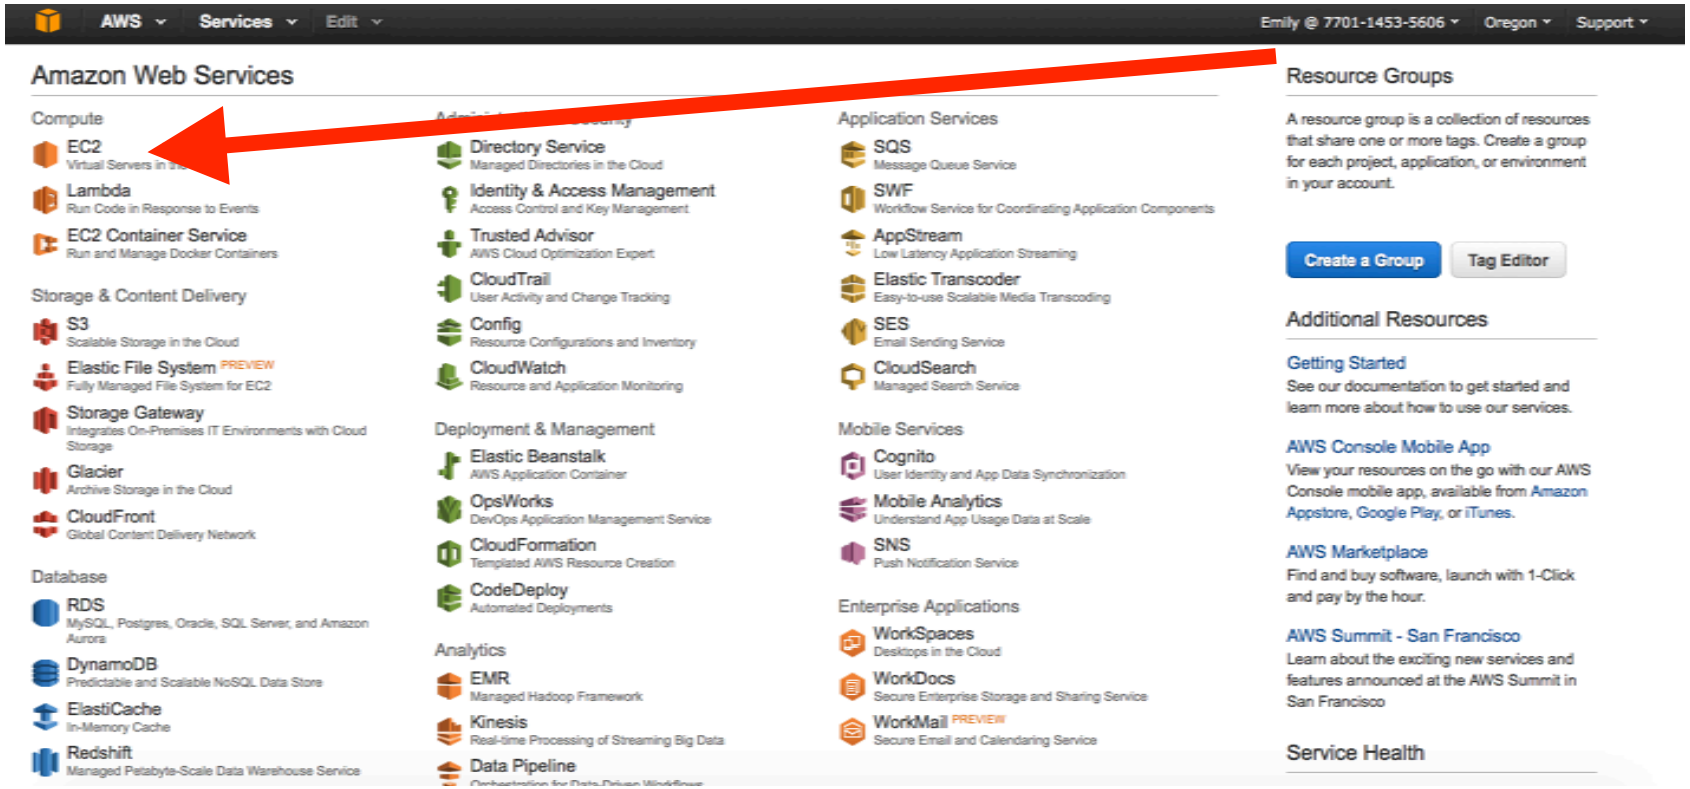

You can do a lot of different things with AWS. For now, let's just start a virtual machine. Click on EC2 ("Elastic Cloud Compute")

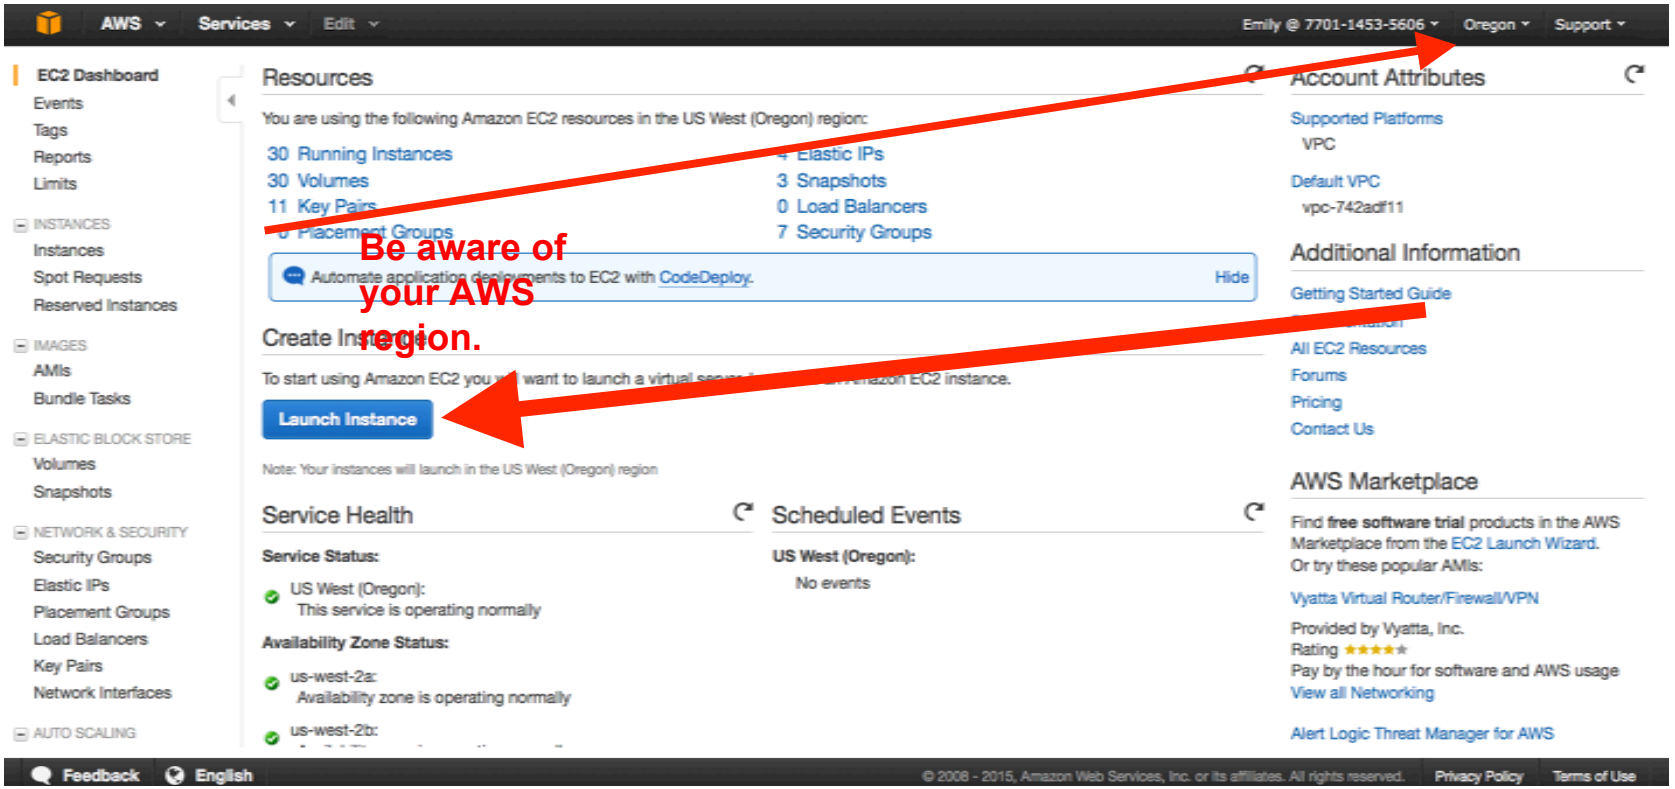

Click on “Launch Instance”. It’s worth noting that AWS has different regions, and that you can launch an instance in any of them. (So if you’ve always wanted a server in Brazil or Japan, this is your chance!) Your choice will affect both latency and price - by default, it’s best to go with Oregon, California, or Virginia.

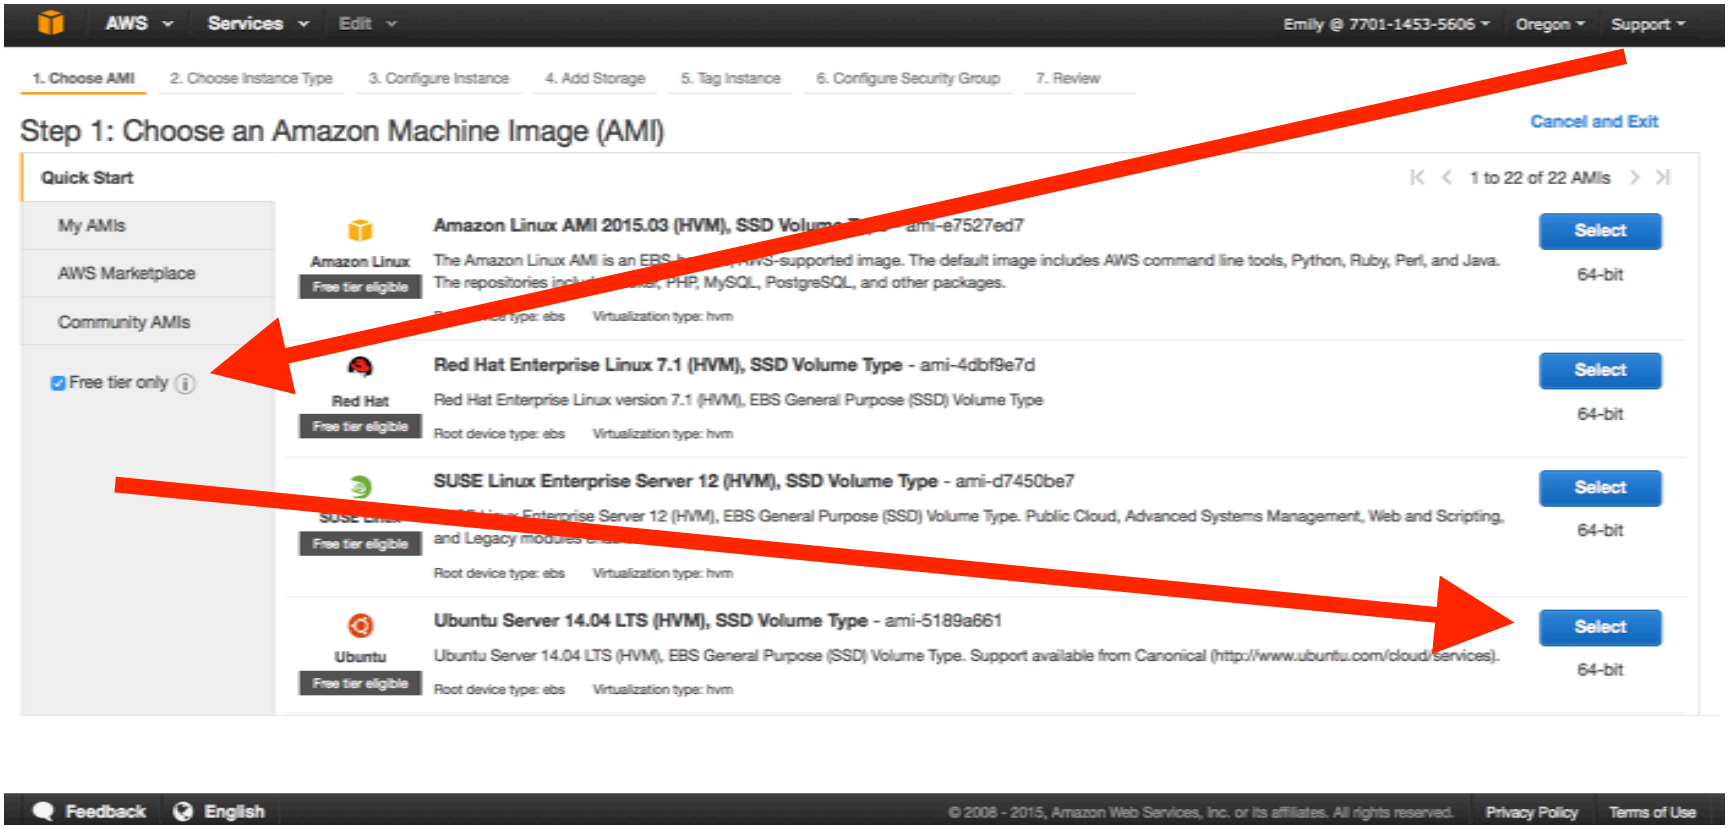

Now it's time to select the image for your VM. For this tutorial, we'll use Ubuntu Server 14.04 LTS (HVM). Be sure to select the 64-bit.
If you are worried about price: filter with the Free tier only box.

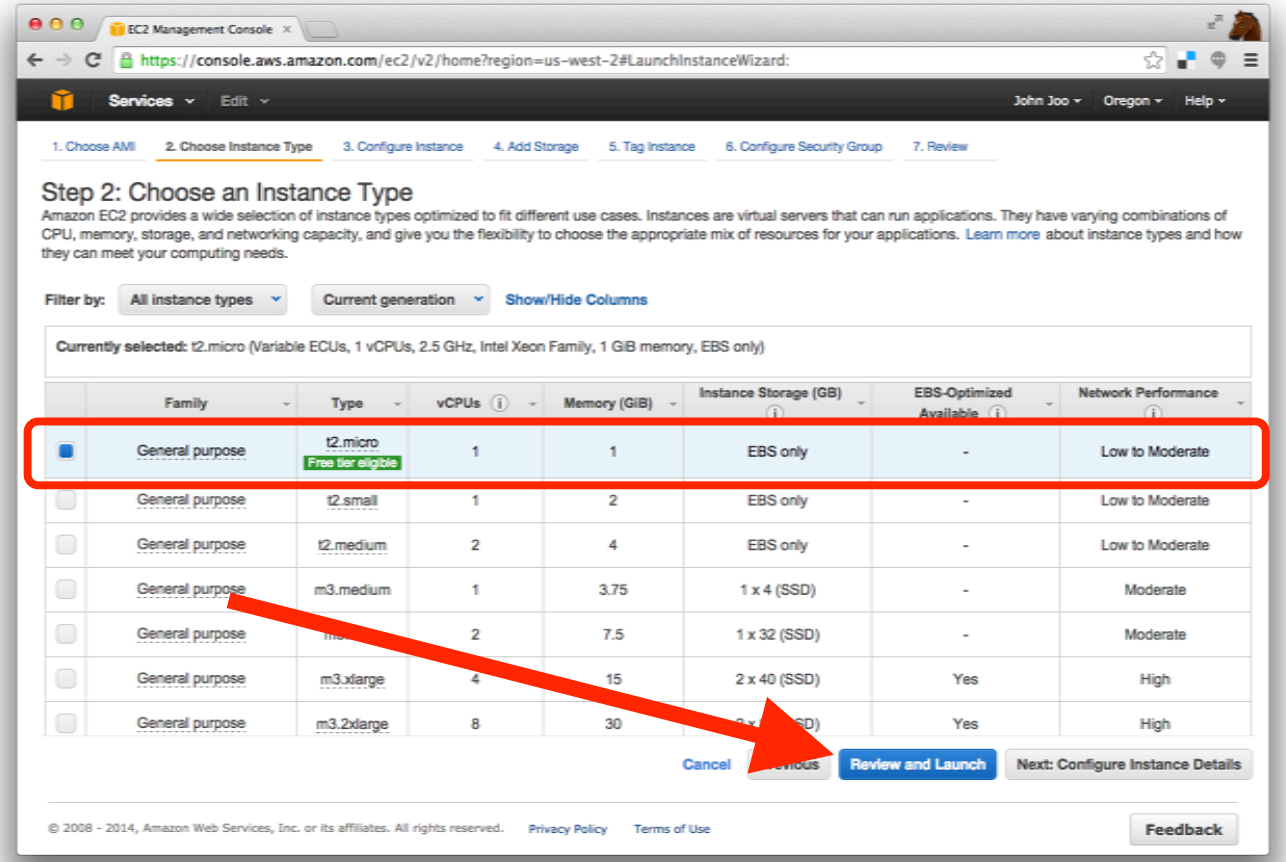

Now choose your instance size. You should use the free option to start: t2.micro. That's what we'll use for this tutorial. (you may need to configure the instance details). Click on "Review and Launch"

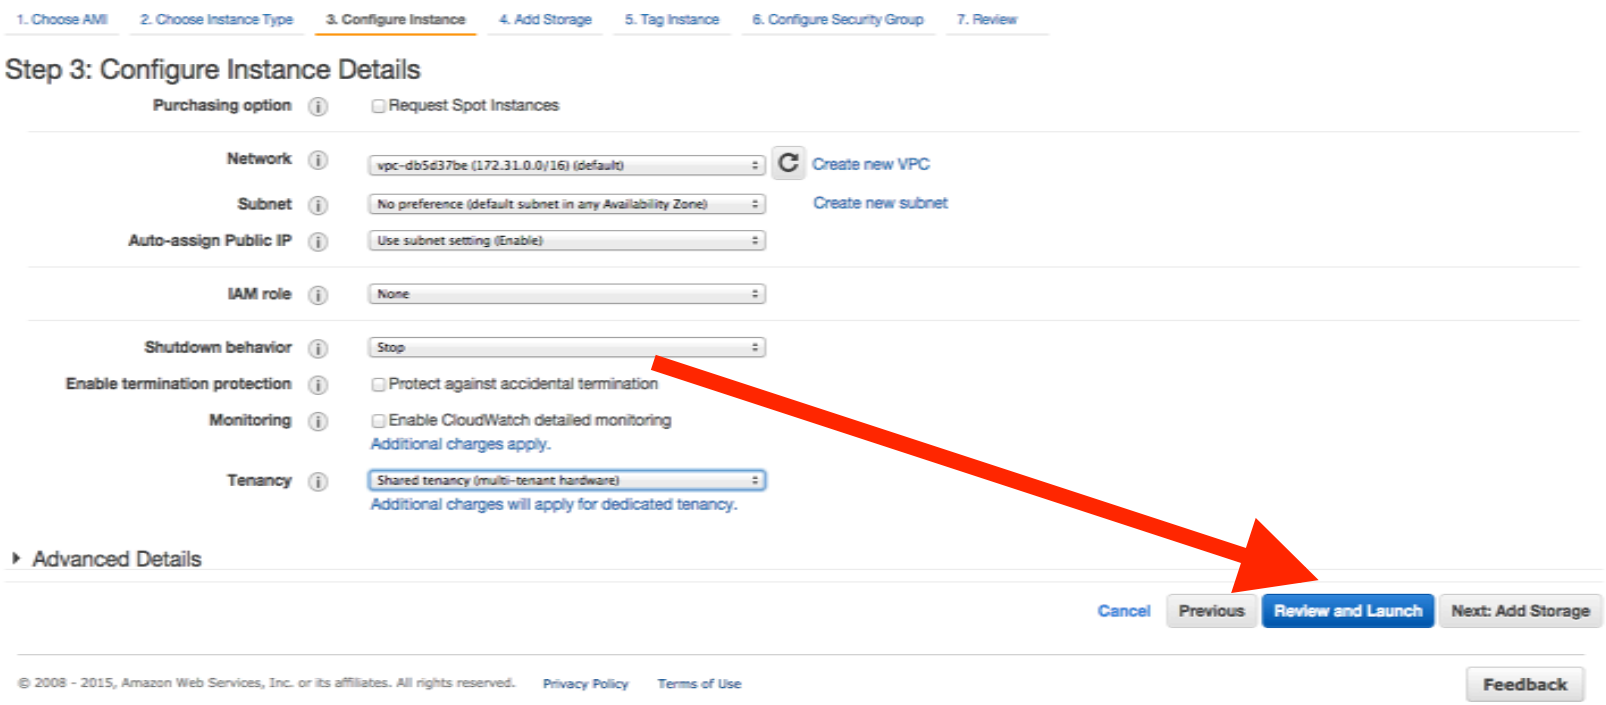

You may need to launch into a VPC (skip this step if this screen doesn't show up)

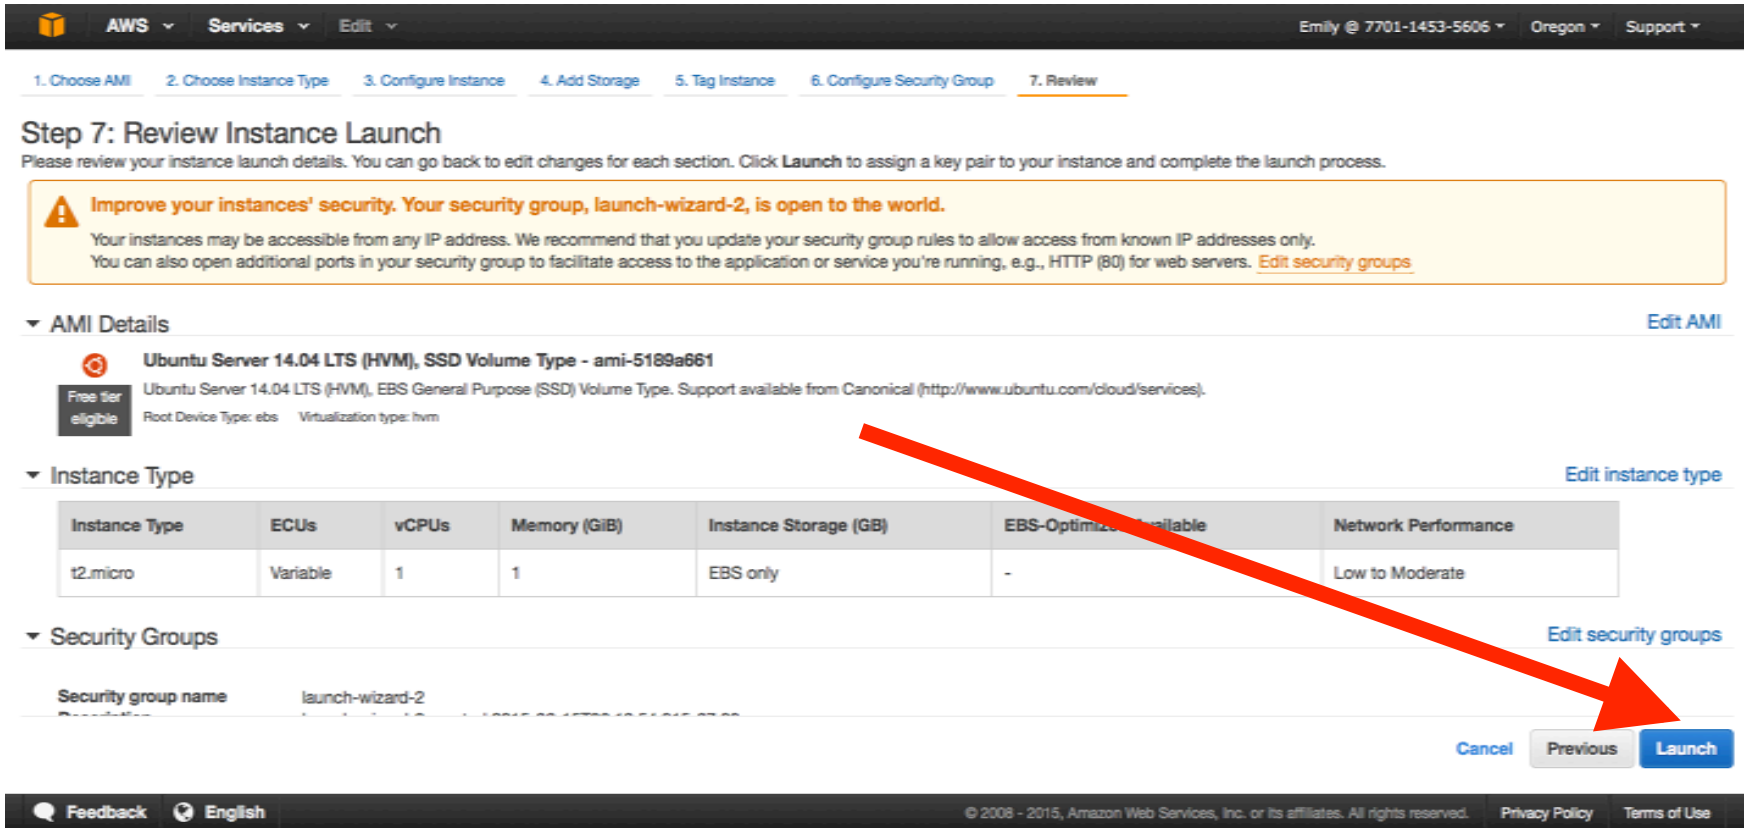
AWS
Services
Edit
Emily @ 7701-1453-5606
Oregon
Support
1. Choose AMI
2. Choose Instance Type
3. Configure Instance
4. Add Storage
5. Tag Instance
6. Configure Security Group
7. Review
Step 7: Review Instance Launch
Please review your instance launch details. You can go back to edit changes for each section. Click Launch to assign a key pair to your instance and complete the launch process.
Improve your instances' security. Your security group, launch-wizard-2, is open to the world.
Your instances may be accessible from any IP address. We recommend that you update your security group rules to allow access from known IP addresses only.
You can also open additional ports in your security group to facilitate access to the application or service you're running, e.g., HTTP (80) for web servers. Edit security groups
AMI Details
Edit AMI
Ubuntu Server 14.04 LTS (HVM), SSD Volume Type - ami-5189a661
Free tier eligible
Ubuntu Server 14.04 LTS (HVM), EBS General Purpose (SSD) Volume Type. Support available from Canonical (http://www.ubuntu.com/cloud/services).
Root Device Type: ebs
Virtualization type: hvm
Instance Type
Edit instance type
Instance Type
ECUs
vCPUs
Memory (GiB)
Instance Storage (GB)
Network Performance
t2.micro
Variable
1
1
EBS only
-
Low to Moderate
Security Groups
Edit security groups
Security group name
launch-wizard-2
Cancel
Previous
Launch
Feedback
English
© 2008 - 2015, Amazon Web Services, Inc. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use

Launch!

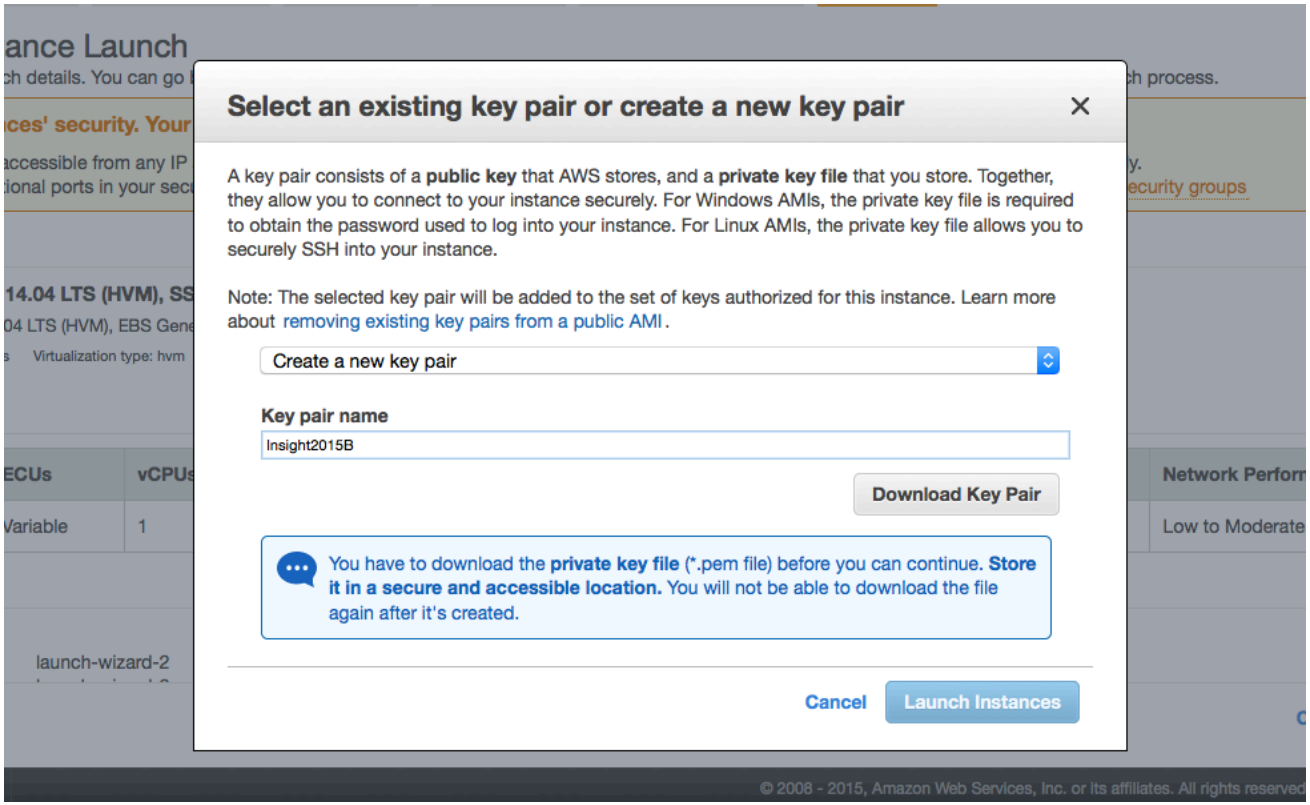

AWS will prompt you to create a key pair for your instance. (You'll need to ssh into the machine). Choose 'Create a new key pair', name it (e.g. 'Insight2015B'), and click 'Download Key Pair'. Then click 'Launch Instances'.

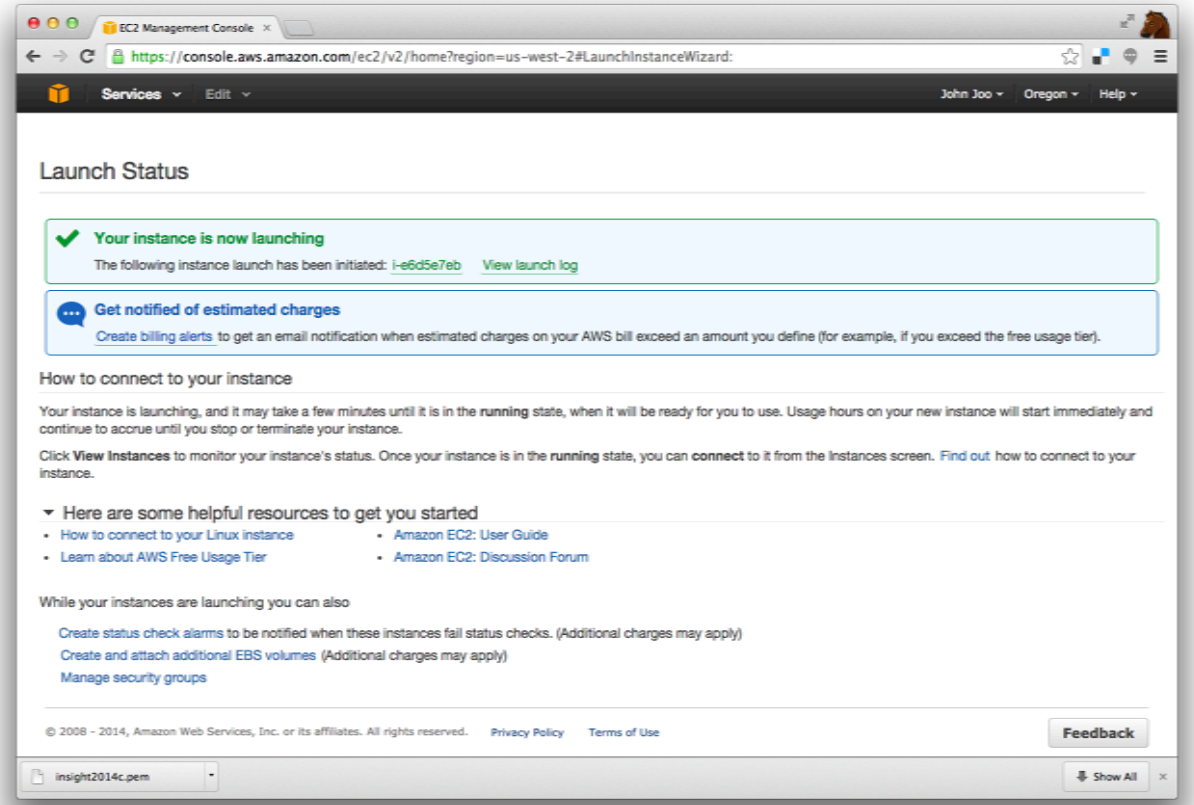

Congratulations! Your AWS instance is now spinning up! Let's log into it. Click on orange box in upper left corner, and then again on EC2

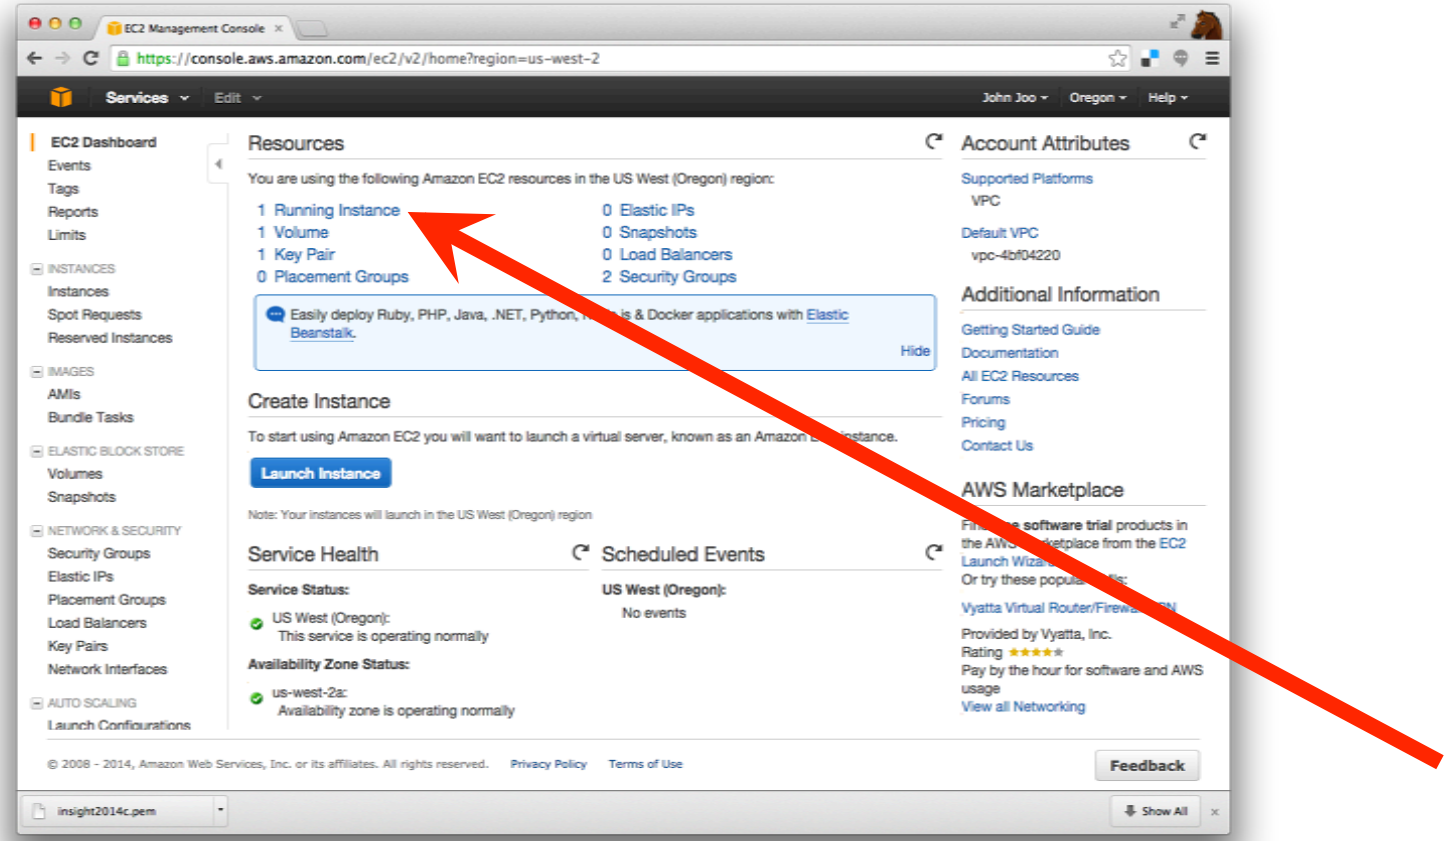

You should now see “1 Running Instance”. Click on it.

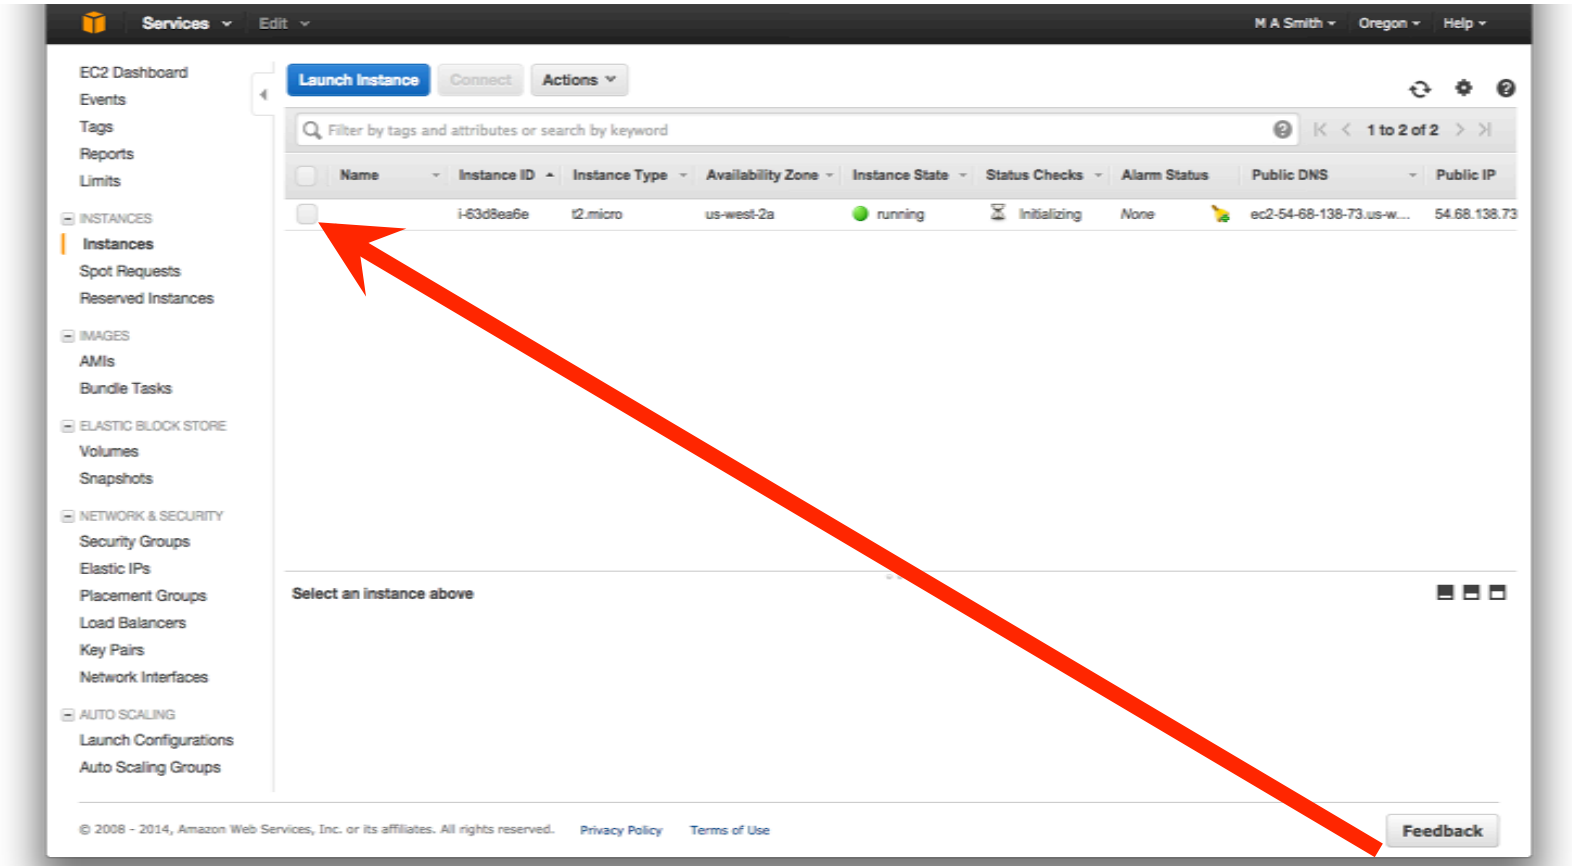

And click again….

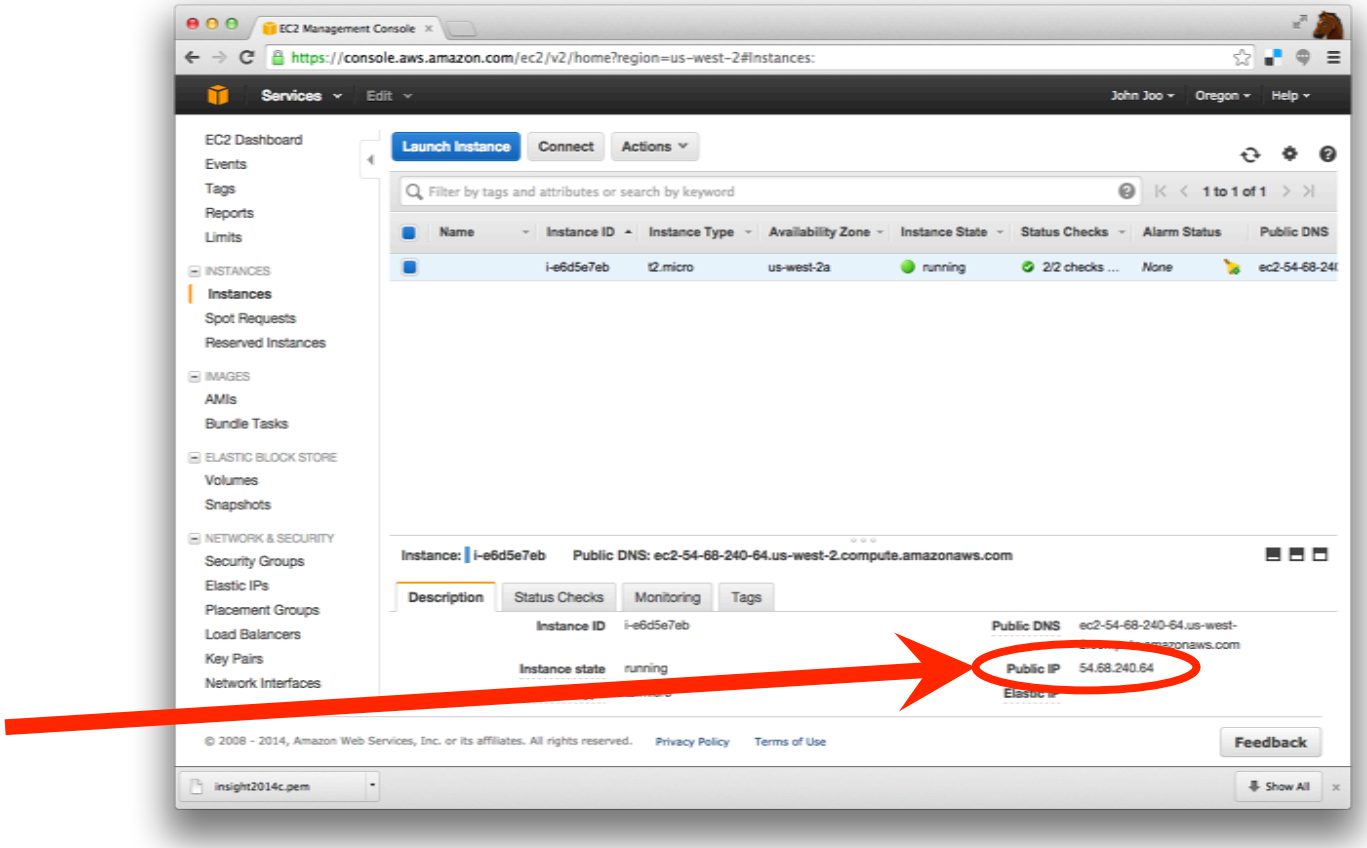

You should now be able to see the public IP address of your VM. This is the endpoint we're going to use to ssh into the machine.

Change permissions.

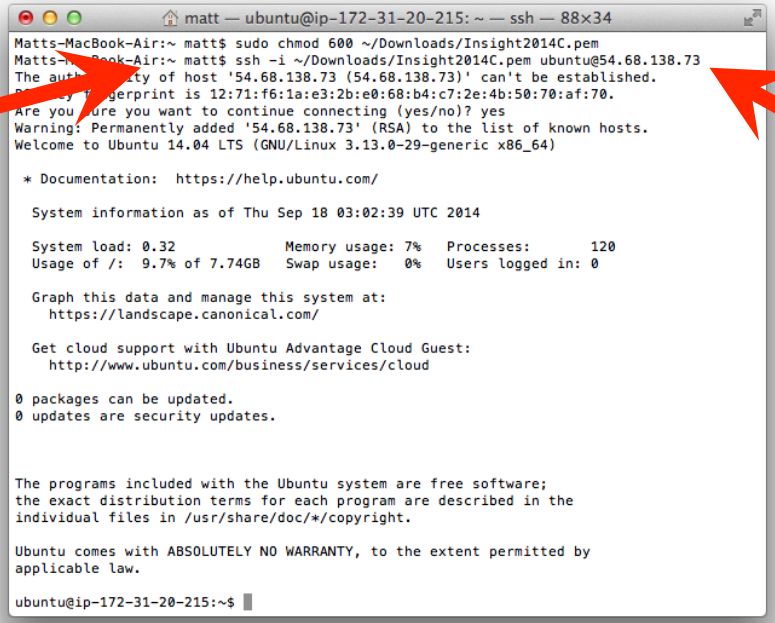

Your IP address

Next….open a terminal!

You need to change the permissions of the private key that you downloaded from AWS earlier (if you didn't put it somewhere special, it's probably in ~/Desktop)

```
> sudo chmod 600 <path to your key>.
```

Now ssh into the machine. Use the command:

```
> ssh -i <path to your key> ubuntu@<your VM's public IP address>
```

If prompted "Are you sure that you want to continue?", enter "yes".

After you're ssh'ed in, let's set up a simple Flask app.
First, we'll need to install a few things. Run the following commands, entering "yes" if prompted.

```
> sudo apt-get update
> sudo apt-get install python-pip python-dev build-essential
> sudo pip install flask
```

Using vim (or pico or emacs or whatever), write the following to the file hello.py:

Now exit and run the co

Intermission: Tips and tricks for getting your SQL database (if you have one) on AWS:

Do you need to install MySQL on your AWS instance? Follow these steps!

Install mysql on your AWS server instance:

```
sudo apt-get install mysql-server
sudo apt-get install mysql-client
```

On your local machine, (or wherever you have your DB initially):

```
mysqldump -u root -p [password] [database name] > dumpFileName.sql
```

note that depending on how you installed mysql, on a mac you might need to specify the path:

```
/usr/local/mysql/bin/mysqldump .. etc..
```

Copy this file to AWS:

```
scp -i <path to your key> dumpFileName.sql ubuntu@<your VM's public IP address>:.
```

in your AWS instance (after creating your [database name]):

```
> mysql
   > create database [database name]
mysql -u root -p [password] [database name] < dumpFileName.sql
```

Unrelated pro tip: In case your server stops running properly:

```
sudo service mysql restart
```

Before you can reach your machine on port 5000, you'll need to open it in your AWS security group settings.

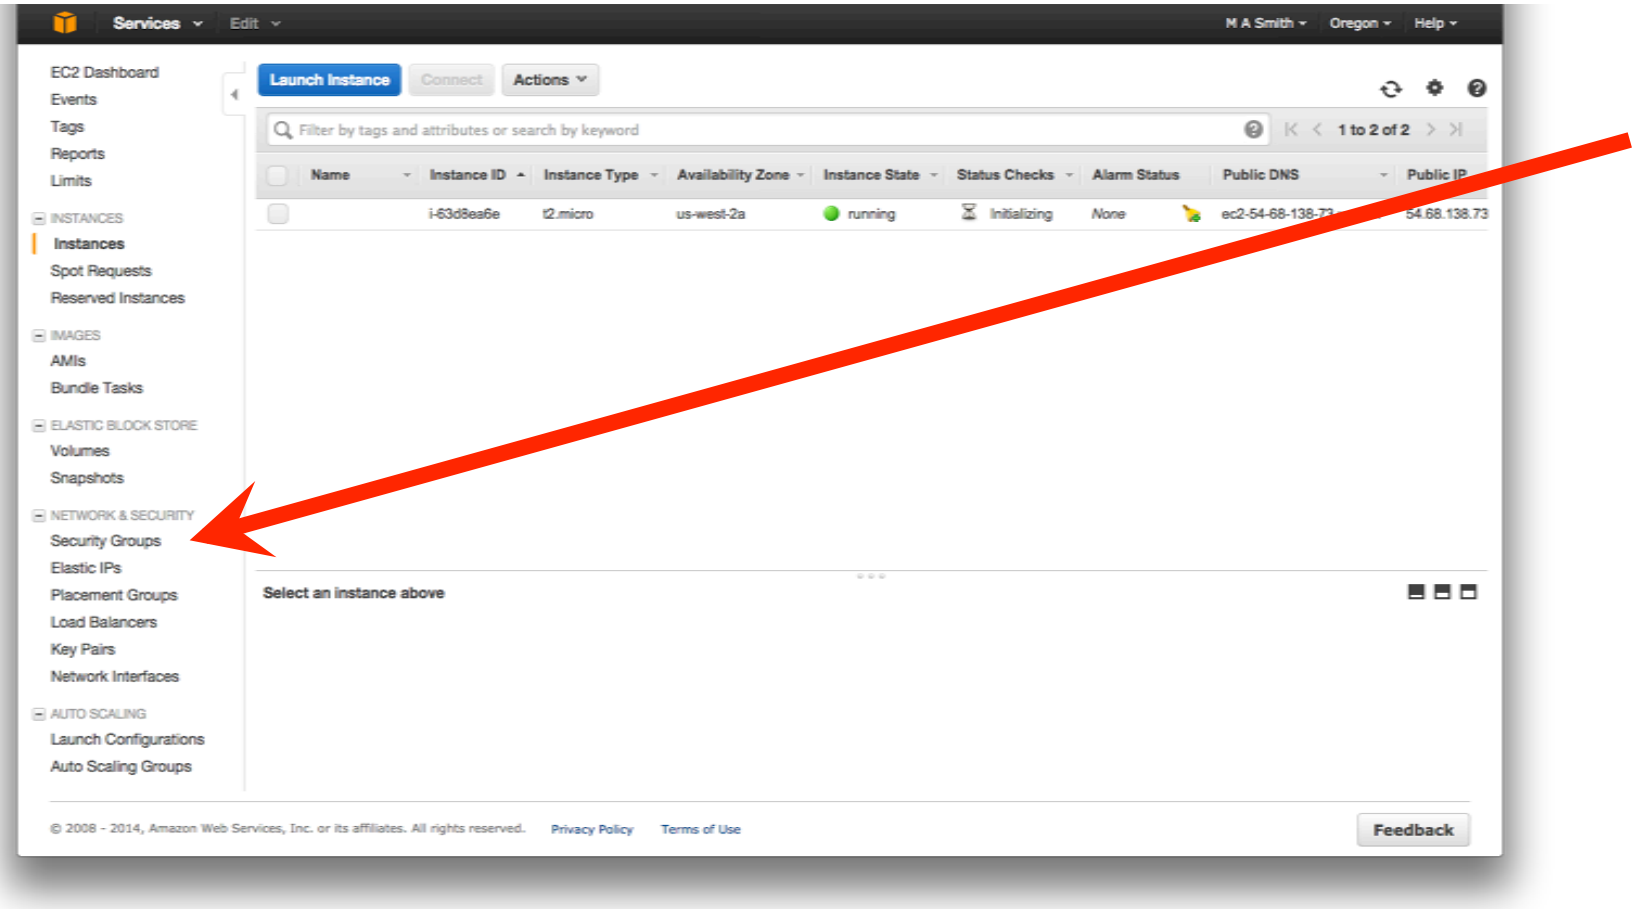

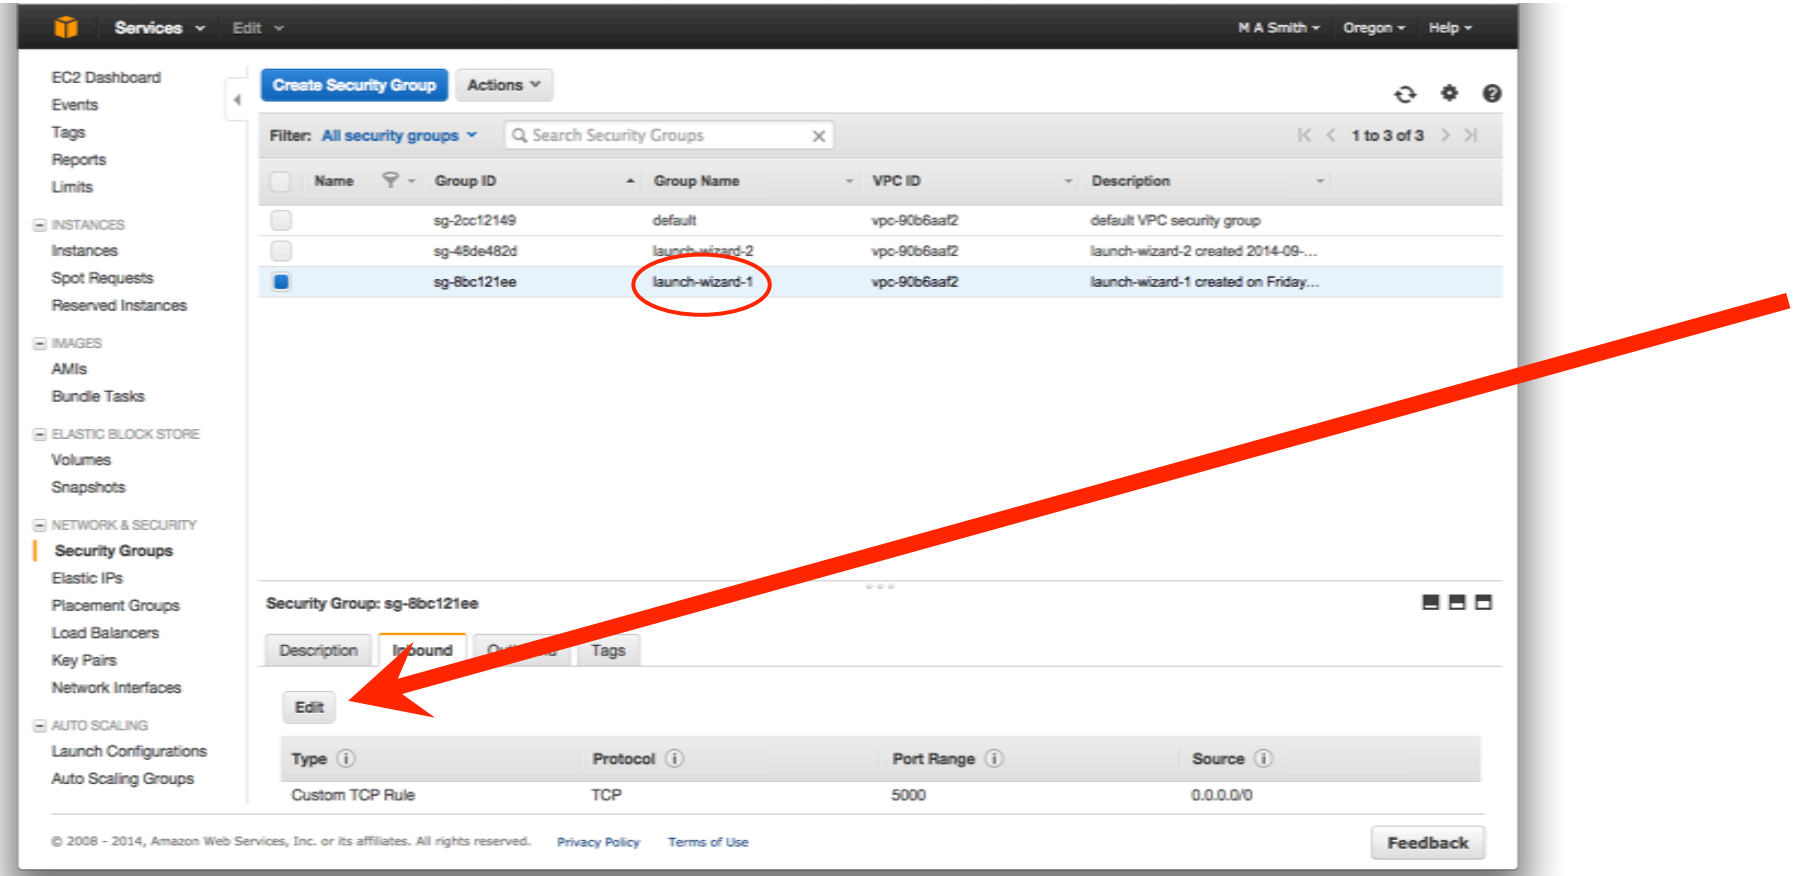

Click Inbound, then Edit

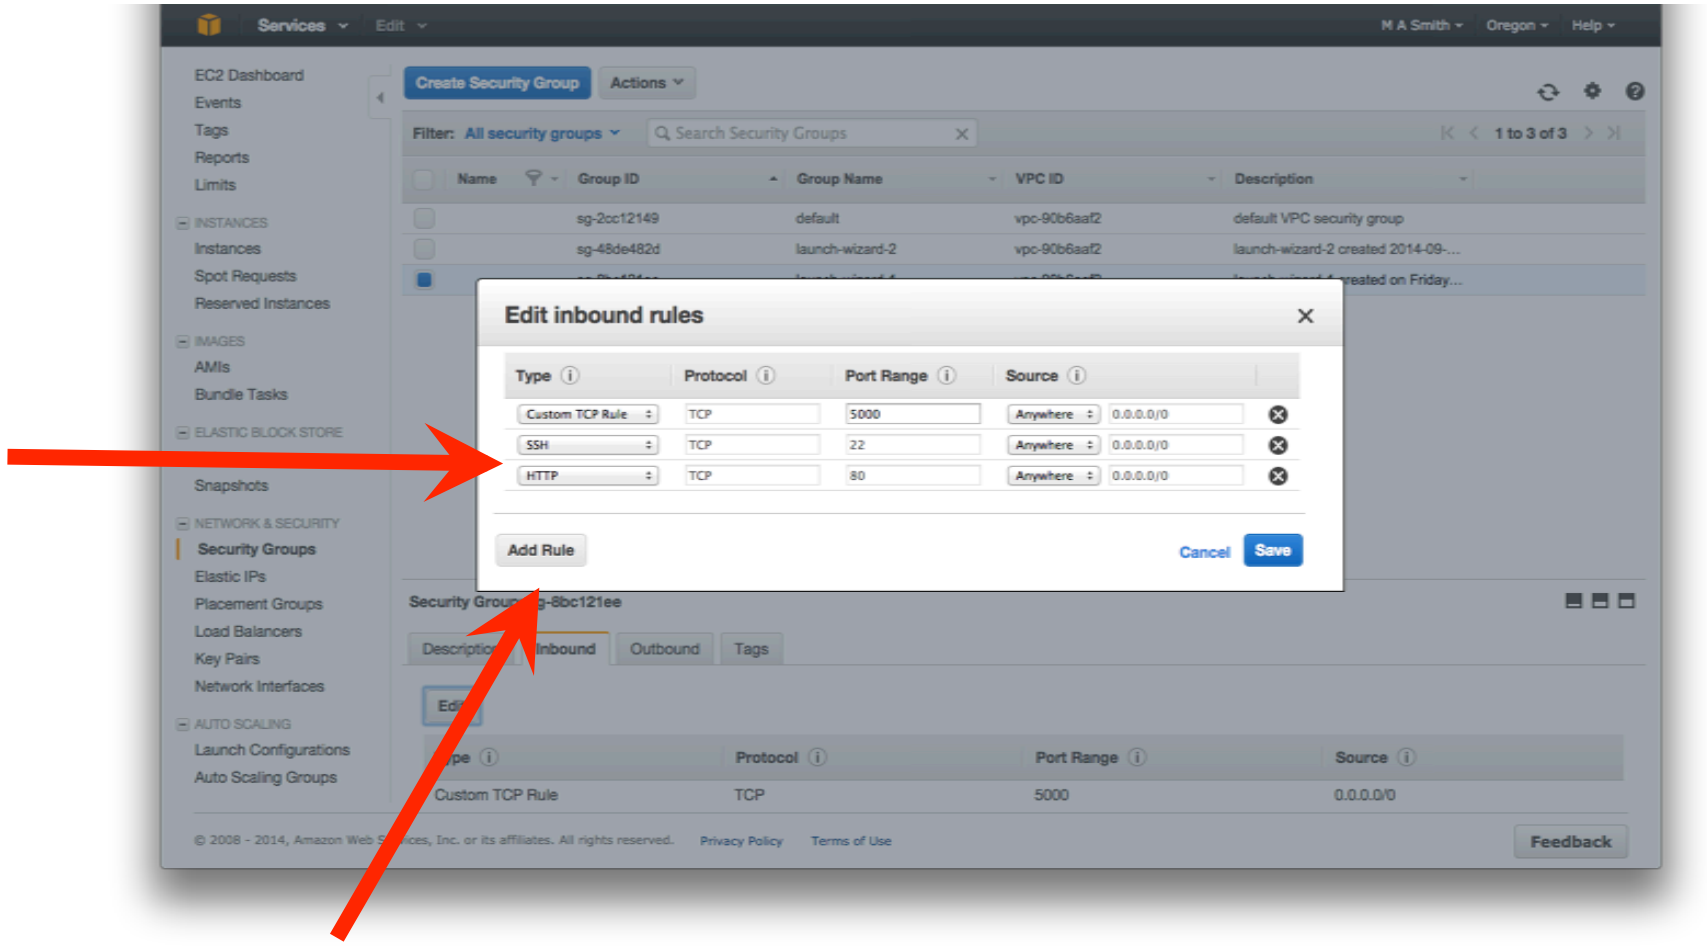

Click Add Rule and add a Custom TCP Rule -> port 5000
Also add SSH, HTTP rules as above

Now use your browser to visit your machine on port 5000.
The url you should enter is: `<the IP address>:5000`

Got it? ur table!

Congratulations!

You now have a running web server online, which anyone in the world (even your mum) can visit!
However, your deployment is not yet production ready.
Some issues:

- If you exit your terminal, the server will die.
- Flask's built-in server isn't production-quali  f too many people visit at once, the server will crash.
- You're on port 5000, rather than 80. 

Bonus Session

The 80% solution: gunicorn + supervisor.

Note: A possible reason that gunicorn is being used here may be that normal users on linux cannot open low ports. However, on AWS you have sudo access, so just using the line in the answer:
http://serverfault.com/questions/112795/how-can-i-run-a-server-on-linux-on-port-80-as-a-normal-user

should work out just fine.

Go back to your security groups on AWS and open port 80.
The Flask webserver is kind of unstable. For a more stable, but slightly more complicated platform, use pip install (or apt-get) to get gunicorn and supervisor
- Create a file called simple.conf with the contents below:

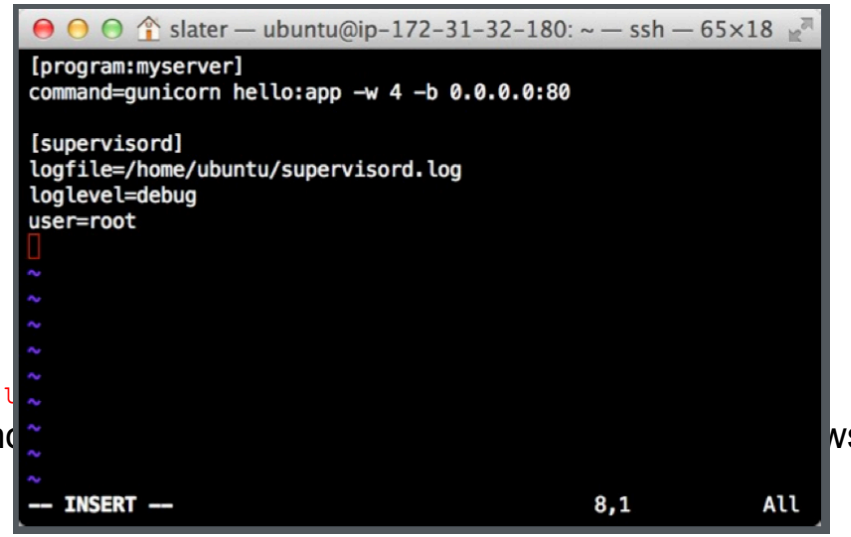

- Run the command su
- Verify that you can nc wser, enter your IP address in the URL bar).

That’s it!

Congratulations on a successful deployment! When deploying your Insight project, please be aware that different settings for gunicorn and supervisor may be more appropriate for your specific situation. The documentation for both projects is very good, so you should be able to figure out what you need. In deciding whether your configuration is “good enough”, there are three important questions to ask:

- Will my project work if 20 people try to use it at the same time? (A good test is to ask other fellows to all visit your site at the same time.)
- If my server crashes, can I quickly (in less than one minute) restart it?
- If my server crashes, are my logs detailed enough for me to figure out what went wrong?